

**ROAD DAMAGE DETECTION USING IMAGE
PROCESSING
A PROJECT REPORT**

Submitted by

AASHIQ MOHAMED R - 61782323103002

BENITA MARY W - 61782323103015

in partial fulfilment for the award of the degree

of

BACHELOR OF ENGINEERING

in

ELECTRONICS AND COMMUNICATION ENGINEERING

**SONA COLLEGE OF TECHNOLOGY, SALEM
(AUTONOMOUS)**

ANNA UNIVERSITY: CHENNAI 600 025

APRIL 2026

BONAFIDE CERTIFICATE

Certified that this project report titled ROAD DAMAGE DETECTION USING IMAGE PROCESSING is the Bonafide work of AASHIQ MOHAMED R (61782323103002) and BENITA MARY W (61782323103015), who carried out the project work under my supervision. The project work was conducted in the Department of Electronics and Communication Engineering, Sona College of Technology, Salem, during the academic year 2025-2026, in partial fulfilment of the requirements for the award of the degree of Bachelor of Engineering in Electronics and Communication Engineering of Anna University, Chennai.

SIGNATURE

Dr. R. S. SABEENIAN,
Professor, Head of the Department,
Department of ECE, Sona College of
Technology, Salem - 636005.

SIGNATURE

Dr. C. M. KALAISELVI,
Assistant Professor, Supervisor,
Department of ECE, Sona College of
Technology, Salem - 636005.

Submitted for the Project Viva-Voce examination held on _____.

Internal Examiner

External Examiner

ACKNOWLEDGEMENT

At this pleasing moment of having successfully completed our main project, we wish to convey our sincere thanks and gratitude to our parents and the College management: our chairman Shri. C. VALLIAPPA, and our vice chairmen Shri. CHOCKO VALLIAPPA and Shri. THYAGU VALLIAPPA who provided all the facilities to us.

We would like to express our sincere thanks and gratitude to our principal, Dr. S. R. R. SENTHILKUMAR for motivating us to do our project and for offering adequate duration for completing our project. His visionary leadership and commitment to academic excellence have been a constant inspiration throughout our academic journey.

We are also immensely grateful and thankful to our Head of the Department Dr. R. S. SABEENIAN, M.E., Ph.D., Professor, Department of ECE, for his constructive suggestion and encouragement during our project. His technical expertise, guidance, and unwavering support have been instrumental in shaping the direction and outcome of this project.

With deep sense of gratitude, we extend our earnest and sincere thanks to our project guide Dr. C. M. KALAISELVI, M.E., Ph.D., Assistant Professor, Department of ECE, for her kind guidance during this project. His invaluable technical insights, patient mentoring, and regular feedback sessions have significantly contributed to the quality of this work.

We also extend our sincere thanks to our Course handling faculty and Project coordinator Dr. R. S. SABEENIAN, M.E., Ph.D., Professor, Department of ECE, for their unwavering support and conduct of project reviews. Their organizational support and constructive evaluations kept our work on track throughout the project cycle.

We express our gratitude to the Teaching and Technical staff of the Department of ECE, Sona College of Technology for their support and cooperation during all stages of project development. The laboratory resources and computing facilities provided by the institution were essential for the successful implementation and testing of the system.

We take this opportunity to extend our deep appreciation to our parents, our family, and friends for all that they meant to us during the crucial times of the completion of the project. Their moral support, encouragement, and patience during late-night work sessions have been a tremendous source of strength and motivation.

AASHIQ MOHAMED R, BENITA MARY W

ABSTRACT

This project presents a comprehensive Road Damage Detection System using Image Processing techniques implemented in MATLAB. The system is designed to identify and classify road surface defects, specifically potholes and cracks, using two complementary approaches: a heuristic image processing pipeline and a deep learning-based custom YOLOv4 object detection model. The primary motivation for this work stems from the critical need to automate road infrastructure monitoring to improve safety and reduce maintenance costs.

The heuristic pipeline employs Region of Interest (ROI) masking with trapezoidal road segmentation, adaptive histogram equalization (CLAHE), adaptive thresholding, and morphological operations to isolate road defects. Detected regions are classified into potholes or cracks based on eccentricity metrics, assigned to lanes, and evaluated for severity using area and contrast factors. The system supports single image and batch image processing modes, making it versatile for various deployment scenarios.

The deep learning approach leverages a custom YOLOv4 architecture trained on XML-annotated pothole datasets. The implementation includes automated anchor box estimation using k-means clustering, pretrained YOLOv4-tiny backbone loading for transfer learning, and a full training pipeline with gradient clipping for stability. The model supports GPU and CPU auto-detection for flexible deployment across different hardware configurations.

TABLE OF CONTENTS

CHAPTER NO	TITLE	PAGE NO
	ABSTRACT	iv
	LIST OF TABLES	vii
	LIST OF FIGURES	viii
	LIST OF ABBREVIATIONS	ix
1.	INTRODUCTION	1
1.1	Objectives of the Project	3
1.2	Scope of the Project	4
1.3	Problem Statement	5
1.4	Significance of the Study	6
1.5	Organization of the Report	8
2.	LITERATURE SURVEY	9
2.1	Classical Image Processing Approaches	10
2.2	Deep Learning Based Approaches	12
2.3	Comparative Analysis of Methods	14
3.	SYSTEM REQUIREMENTS	15
3.1	Hardware Requirements	16
3.2	Software Requirements	17
3.3	Dataset Requirements	18
4.	METHODOLOGY	19
4.1	Overall System Architecture	20
4.2	Heuristic Image Processing Pipeline	21
4.3	YOLOv4 Deep Learning Pipeline	24
4.4	Integration and Decision Framework	26
5.	IMPLEMENTATION	27
5.1	Heuristic Detection System Implementation	28
5.2	YOLOv4 Implementation	33

5.3	Visualization and Result Display	38
5.4	Testing and Validation Strategy	39
6.	RESULTS AND DISCUSSION	40
6.1	Heuristic Pipeline Results	41
6.2	Severity Classification Results	43
6.3	YOLOv4 Detection Results	45
6.4	Discussion and Comparative Analysis	47
7.	APPLICATIONS AND FUTURE SCOPE	48
7.1	Applications	49
7.2	Future Scope	51
8.	CONCLUSION	54
9.	REFERENCES	57

LIST OF TABLES

TABLE NO	TITLE	PAGE NO
1	SYSTEM REQUIREMENTS SUMMARY	17
2	COMPARISON OF HEURISTIC VS YOLOV4 DETECTION APPROACHES	28
3	SEVERITY CLASSIFICATION TABLE	43
4	DETECTION RESULTS SUMMARY	45
5	YOLOV4 MODEL PERFORMANCE METRICS	46
6	HARDWARE CONFIGURATION TEST RESULTS	47
7	ABBREVIATIONS AND MEANINGS	ix

LIST OF FIGURES

FIGURE NO	TITLE	PAGE NO
1.1	BLOCK DIAGRAM OF THE ROAD DAMAGE DETECTION SYSTEM	5
2.1	OVERVIEW OF CLASSICAL VS. DEEP LEARNING DETECTION APPROACHES	12
4.1	SYSTEM WORKFLOW / FLOWCHART	21
5.1	ROI MASKING TRAPEZOID GEOMETRY	29
5.2	MORPHOLOGICAL PROCESSING STAGES	32
5.3	YOLOV4 TRAINING PROGRESS PLOT	37
6.1	ANNOTATED DETECTION OUTPUT - SINGLE IMAGE	41
6.2	SEVERITY MAP VISUALIZATION	43
6.3	BATCH PROCESSING RESULTS - MULTIPLE IMAGES	45

LIST OF ABBREVIATIONS

Abbreviation	Meaning
MATLAB	Matrix Laboratory
YOLOv4	You Only Look Once version 4
ROI	Region of Interest
CLAHE	Contrast Limited Adaptive Histogram Equalization
XML	Extensible Markup Language
CNN	Convolutional Neural Network
GPU	Graphics Processing Unit
CPU	Central Processing Unit
IoU	Intersection over Union
bw	Binary (Black and White) Image
SGD	Stochastic Gradient Descent
LED	Light Emitting Diode
GPS	Global Positioning System
RGB	Red, Green Blue

CHAPTER 1

INTRODUCTION

Road infrastructure is among the most critical components of any country's transportation network. The quality of roads directly affects vehicle safety, fuel efficiency, travel time, and the overall economy. Potholes, cracks, and surface deterioration are among the most common forms of road damage, posing significant risks to vehicle occupants and pedestrians alike. Traditional road inspection methods rely heavily on manual visual surveys by maintenance crews, which are time-consuming, labour-intensive, inconsistent, and unable to scale effectively across large road networks.

This project presents a comprehensive Road Damage Detection System implemented using MATLAB. The system employs two complementary approaches. The first is a classical image processing heuristic pipeline that uses region of interest masking, adaptive contrast enhancement, binary thresholding, and morphological operations to detect and classify road surface defects. The second is a deep learning approach based on a custom YOLOv4 (You Only Look Once version 4) object detection model trained specifically on annotated pothole datasets.

1.1 OBJECTIVES OF THE PROJECT

The primary objective of this project is to develop an automated, reliable, and computationally efficient system for road damage detection using image processing and deep learning techniques. The specific objectives are enumerated as follows:

- To develop an automated system for detecting potholes and cracks in road surface images captured by standard cameras under varied real-world conditions.
- To implement a heuristic image processing pipeline using MATLAB incorporating ROI masking, adaptive contrast enhancement (CLAHE), adaptive thresholding, and morphological analysis for defect extraction.
- To design and train a custom YOLOv4 deep learning model for pothole detection using Pascal VOC XML-annotated datasets, incorporating transfer learning from pretrained backbones.

1.2 SCOPE OF THE PROJECT

The scope of this project encompasses the design, implementation, and validation of a MATLAB-based road damage detection system. The system is applicable to images obtained from dashboard cameras, mobile mapping vehicles, or standard digital cameras mounted at road level. It covers defect detection for potholes and surface cracks under various weather conditions including dry, wet, and night scenarios, with adaptive processing adjustments for each condition.

The scope explicitly excludes real-time video processing on embedded hardware platforms, which is deferred to future work. Similarly, integration with live GPS navigation systems, cloud platforms, and IoT frameworks, while discussed as future enhancements, are not implemented within the current project scope. The system focuses on the core detection and classification algorithms, establishing a solid foundation for these extended capabilities.

1.3 PROBLEM STATEMENT

Road maintenance agencies face a persistent challenge: how to efficiently and accurately monitor the condition of road surfaces across extensive networks with limited human resources and budget. Manual inspection by maintenance crews is subjective, inconsistent, and prohibitively expensive at scale. Inspection frequency is often insufficient to detect emerging defects before they escalate into safety hazards or require costly structural repairs.

Deep learning approaches address many of these challenges but require substantial annotated training data and computational resources for model training. Classical image processing methods are faster and more interpretable but less robust under challenging conditions. There exists a need for a dual-approach system that combines the strengths of both methodologies, providing reliable detection across diverse scenarios while remaining deployable on standard hardware.

1.4 SIGNIFICANCE OF THE STUDY

The significance of this study lies in its practical contribution to infrastructure management and public safety. Road surface deterioration, if left undetected and unaddressed, can lead to vehicle damage, traffic accidents, and increased travel times. In India, road accident statistics consistently identify poor road conditions as a significant contributing factor in a substantial proportion of road fatalities and injuries each year.

By providing an automated, scalable, and cost-effective solution for road damage detection, this project contributes directly to public safety improvement. Early detection of critical potholes and cracks enables proactive maintenance scheduling, preventing further structural deterioration and reducing long-term repair costs. This has direct economic significance for road maintenance authorities and road users alike.

1.5 ORGANIZATION OF THE REPORT

The remainder of this report is organized as follows:

Chapter 2 presents a comprehensive literature survey covering existing research on road damage detection using both classical image processing and deep learning approaches. Key methodologies, datasets, and performance benchmarks reported in the literature are reviewed and critically analysed.

Chapter 3 describes the system requirements including hardware specifications, software environment, and dataset requirements for the Road Damage Detection System.

Chapter 4 details the methodology adopted for the system, covering the overall dual-pipeline architecture, the heuristic image processing pipeline, and the YOLOv4 deep learning pipeline.

Chapter 5 presents the implementation details of both pipelines, including code structure, key algorithmic implementations, and visualization components.

Chapter 6 presents the results and discussion, including quantitative and qualitative evaluation of both detection approaches, severity classification analysis, and comparative performance discussion.

Chapter 7 discusses the applications of the system across various domains and outlines the future scope for system enhancement and extension.

Chapter 8 provides the conclusion summarizing the key contributions, findings, and overall outcomes of the project.

The report concludes with a comprehensive list of references cited throughout the document.

CHAPTER 2

LITERATURE SURVEY


Research in automated road condition assessment has grown substantially over the past two decades. The proliferation of low-cost imaging hardware, advances in computer vision algorithms, and the emergence of deep learning have collectively enabled increasingly sophisticated approaches to road surface defect detection. This chapter reviews the major categories of approaches reported in the literature, highlighting key contributions, methodological innovations, and performance benchmarks.

2.1 CLASSICAL IMAGE PROCESSING APPROACHES

Early approaches to automated road damage detection relied on specialized vehicle-mounted sensors such as laser profilometers, ground-penetrating radar, and acoustic sensors. While highly accurate, these methods are expensive and require dedicated inspection vehicles, making large-scale deployment impractical for most road maintenance organizations.

The introduction of low-cost cameras and the proliferation of smartphones enabled image-based road inspection approaches. Among the earliest image processing methods, researchers used edge detection operators such as Sobel, Canny, and Laplacian of Gaussian to identify road surface discontinuities. These gradient-based approaches detected sharp transitions in intensity that correlate with crack boundaries. However, they were highly sensitive to noise, shadows, and lighting variations, producing many false positives on textured surfaces.

Figure 1.1: Block Diagram of the Road Damage Detection System



2.2 DEEP LEARNING BASED APPROACHES

The introduction of deep learning, particularly Convolutional Neural Networks (CNNs), marked a paradigm shift in road damage detection. ResNet, VGG, and MobileNet architectures were adapted for road surface classification. The YOLO (You Only Look Once) family of object detectors gained popularity for real-time detection, with YOLOv3 and YOLOv4 achieving excellent trade-offs between speed and accuracy. Researchers demonstrated that YOLOv4, with its CSPDarknet53 backbone, SPP module, and PANet neck, provided superior multi-scale detection capabilities suitable for detecting potholes of varying sizes.

2.3 WEATHER ADAPTATION AND PREPROCESSING

Weather condition adaptation is an area of active research in road damage detection. Night-time and wet condition images pose unique challenges due to reflections, reduced contrast, and altered road surface appearance. Studies have proposed pre-processing pipelines incorporating CLAHE, gamma correction, and histogram stretching to normalize images before defect detection algorithms are applied.

Gopalakrishnan et al. (2017) demonstrated that transfer learning from ImageNet-pretrained deep convolutional networks significantly improved pavement distress detection accuracy, particularly for small training datasets typical of specialized road damage annotation projects. Their work validated the effectiveness of transfer learning

The review of existing literature reveals several research gaps that this project addresses: (1) the lack of dual-approach systems that combine classical and deep learning methods within a single framework; (2) limited work on severity classification using composite scoring from single-camera imagery; (3) insufficient focus on weather-adaptive processing for road damage detection; and (4) the need for accessible, MATLAB-based implementations that support academic research and educational applications.

CHAPTER 3

SYSTEM REQUIREMENTS

The development and implementation of the Road Damage Detection System require a combination of software, computational hardware, and data resources. The system is implemented entirely in MATLAB and can operate on standard personal computers, though GPU acceleration is recommended for the YOLOv4 deep learning component. This chapter details the hardware, software, and dataset requirements for deploying and operating the system.

3.1 HARDWARE REQUIREMENTS

3.1.1 COMPUTATIONAL PLATFORM

The computational requirements for the Road Damage Detection System vary depending on the detection mode used. The heuristic pipeline can be executed efficiently on standard personal computers with moderate specifications. The YOLOv4 deep learning pipeline benefits significantly from GPU acceleration, particularly during the training phase. The following table summarizes the minimum and recommended hardware specifications.

Parameter	Minimum Requirement	Recommended
Processor	Intel Core i5 or equivalent	Intel Core i7 / AMD Ryzen 7
RAM	8 GB	16 GB or higher
GPU (for YOLOv4)	Not required (CPU mode)	NVIDIA GTX 1060 or higher (CUDA)
Storage	50 GB free space	100 GB SSD
Display	1280 x 720 resolution	1920 x 1080 or higher

Table 1: System Requirements Summary

3.1.2 CAMERA / IMAGE ACQUISITION DEVICE

The quality of road surface images significantly impacts detection performance. The following camera specifications are recommended for optimal system performance:

- Dashboard camera capable of capturing at minimum 720p resolution images with a frame rate of at least 30 fps for video processing applications.
- Standard digital camera or smartphone with at least 8-megapixel sensor for collecting training and test images.
- UAV or drone camera for aerial road surface inspection, enabling wider coverage and overhead perspectives for large-scale road monitoring deployments.
- Industrial-grade road inspection cameras with vibration dampening for vehicle-mounted deployment in active road survey operations.

3.2 SOFTWARE REQUIREMENTS

3.2.1 MATLAB ENVIRONMENT

The entire system is developed using MATLAB R2021a or later versions. MATLAB provides a comprehensive environment for image processing, machine learning, and deep learning with extensive toolbox support. The following MATLAB toolboxes are required for full system functionality:

- Image Processing Toolbox - for adaptive thresholding, morphological operations, region property analysis, and image annotation functions including binarize, regionprops, adapthisteq, and poly2mask.
- Computer Vision Toolbox - for YOLOv4 object detector creation using yolov4ObjectDetector, training configuration, and inference functions. Required for anchor box estimation using estimateAnchorBoxes.
- Deep Learning Toolbox - for neural network layer definition, training using trainNetwork, and network creation for custom architecture deployment.
- Parallel Computing Toolbox (optional) - for GPU-accelerated YOLOv4 training, enabling significant reduction in training time when CUDA-compatible GPUs are available.

Software Component	Version / Details
MATLAB	R2021a or later
Image Processing Toolbox	Required
Computer Vision Toolbox	Required for YOLOv4
Deep Learning Toolbox	Required for YOLOv4
Operating System	Windows 10/11, Ubuntu 20.04+
CUDA Toolkit (optional)	11.0 or higher for GPU support

3.3 DATASET REQUIREMENTS

The Road Damage Dataset (RDD) is a publicly available benchmark dataset that provides annotated road damage images from multiple countries. The RDD dataset includes annotations for multiple damage categories including longitudinal cracks, transverse cracks, alligator cracks, and potholes, making it suitable for training the YOLOv4 model in this project. Alternatively, custom datasets annotated using tools such as Labelling or CVAT can be prepared and used with the system.

Image quality requirements for training data include minimum resolution of 640x480 pixels, JPEG or PNG format, and clear visibility of road surface defects with minimal motion blur. A balanced distribution of defect sizes across small, medium, and large categories improves model generalization. Data augmentation strategies including horizontal flipping, brightness variation, and geometric transforms are applied during training to increase effective dataset size and diversity.

CHAPTER 4


METHODOLOGY

The methodology adopted for the Road Damage Detection System follows a dual-pipeline approach, combining classical image processing techniques with deep learning-based object detection. The system is structured to provide reliable detection across varied road conditions, with each pipeline contributing distinct capabilities to the overall system. This chapter describes the architectural design, algorithmic choices, and processing strategies employed in both pipelines.

4.1 OVERALL SYSTEM ARCHITECTURE

The system accepts road surface images in single or batch format. Upon input, the user specifies the weather condition (Dry, Wet, or Night) and optional GPS coordinates for geospatial logging. The heuristic pipeline then performs ROI extraction, image enhancement, binary segmentation, morphological filtering, and region property analysis to identify and classify defects. Independently, the YOLOv4 pipeline loads or trains a detection model, runs inference on input images, and returns bounding boxes with confidence scores. The outputs of both pipelines are combined in a unified visualization and reporting module.

Figure 2.1: Overview of Classical vs. Deep Learning Approaches in Road Damage Detection



4.2 HEURISTIC IMAGE PROCESSING PIPELINE

4.2.1 REGION OF INTEREST (ROI) MASKING

The first step of the heuristic pipeline applies a trapezoidal ROI mask to focus processing exclusively on the road surface area. The trapezoid is defined by four vertices: bottom-left at 5% of image width, mid-left at 40% width and 45% height, mid-right at 60% width and 45% height, and bottom-right at 95% width. This geometry mimics the visible road area from a typical dashboard camera perspective. Non-road pixels are set to white (255), effectively eliminating buildings, sky, and vegetation from subsequent processing.

The trapezoidal ROI design is motivated by the perspective geometry of road surface images captured from a forward-facing camera mounted at vehicle height. The road surface appears in the lower portion of the image as a trapezoidal region, narrowing towards the horizon. By masking out the upper and peripheral regions, the pipeline

eliminates most non-road pixels that would otherwise generate false detections from sky textures, building edges, and vegetation boundaries.

4.2.2 CONTRAST ENHANCEMENT

After ROI masking, the grayscale road surface image undergoes Contrast Limited Adaptive Histogram Equalization (CLAHE) using MATLAB's `adapthisteq` function with a clip limit of 0.03. CLAHE enhances local contrast while preventing over-amplification of noise in near-uniform regions. This is particularly important for wet road surfaces where specular reflections can reduce the visibility of defect boundaries.

Unlike global histogram equalization, which applies a single transformation to the entire image, CLAHE operates on local tiles of the image. Each tile undergoes independent histogram equalization, with the clip limit controlling the maximum slope of the cumulative distribution function to prevent noise amplification. The enhanced tiles are then combined using bilinear interpolation to produce a smooth output image without tile boundary artifacts.

4.2.3 ADAPTIVE THRESHOLDING AND MORPHOLOGICAL PROCESSING

The contrast-enhanced image is converted to a binary image using adaptive thresholding with a sensitivity parameter of 0.45. Adaptive thresholding computes local threshold values based on neighbourhood statistics, making it robust to illumination gradients. The binary image is then inverted so that dark defect regions appear as white foreground objects. Morphological opening with a disk-shaped structuring element of radius 2 removes small noise pixels, while morphological closing with radius 4 bridges small gaps within defect regions.

4.2.4 REGION PROPERTY ANALYSIS AND CLASSIFICATION

Connected component analysis using region props extracts bounding boxes, area, solidity, and eccentricity for each detected blob. Filtering criteria require detected regions to have area greater than 250 pixels and solidity greater than 0.5, ensuring that only compact, substantial defects are retained. Eccentricity is used for type

classification: regions with eccentricity greater than 0.92 are classified as Cracks (elongated shapes), while those with lower eccentricity are classified as Potholes (roughly circular shapes). Each detection is then assigned to a lane based on its horizontal centre position.

4.2.5 SEVERITY CALCULATION

Severity is computed as a composite score combining a size factor (proportional to bounding box area, capped at 50) and a contrast factor (based on mean patch intensity relative to maximum, scaled to 50). Under wet conditions, the severity score is multiplied by 1.2 to account for increased hazard. The final severity score is bounded between 0 and 100. Risk levels are assigned as: Critical (severity > 70), Medium (severity 35-70), and Low (severity < 35).

The size factor captures the physical extent of the defect, which is directly related to the risk it poses to vehicles. Larger potholes and longer cracks pose greater risks and require more urgent repair attention. The contrast factor captures the depth indication of the defect: darker regions represent deeper depressions with stronger shadowing, indicating more severe damage. By combining both factors, the severity score provides a more comprehensive assessment than either factor alone.

4.3 YOLOV4 DEEP LEARNING PIPELINE

4.3.1 DATASET PREPARATION

The YOLOv4 training pipeline begins with parsing XML annotation files from the training directory. A MATLAB container. Map indexes all image files by filename for rapid lookup. The parse XML Annotations function isolates DOM object creation within a separate function scope to prevent MATLAB workspace serialization issues. Valid bounding boxes with width and height exceeding 5 pixels are retained and stored in a MATLAB table structure compatible with image Datastore and box Label Datastore.

Data augmentation strategies applied during training include random horizontal flipping, random brightness adjustment within $\pm 20\%$ of the original value, random

contrast adjustment, and random crop and resize operations to simulate different camera distances and angles. These augmentations increase the effective diversity of the training dataset without requiring additional manual annotation, improving the model's generalization capability.

4.3.2 ANCHOR BOX ESTIMATION

Nine anchor boxes are estimated from the training dataset using MATLAB's estimate Anchor Boxes function, which employs k-means clustering on bounding box dimensions. The nine anchors are split into three groups of three for use at the three detection scales of YOLOv4: small-scale anchors for large objects, medium-scale anchors for medium objects, and large-scale anchors for small objects. Mean IoU of the estimated anchors is reported to assess clustering quality.

4.3.3 MODEL INITIALIZATION AND TRAINING

The YOLOv4 detector is initialized preferentially using a pretrained YOLOv4-tiny backbone, which provides transfer learning benefits and significantly reduces training time. If the pretrained model is unavailable, a custom lightweight YOLOv4-like architecture is constructed with Conv-BN-ReLU blocks forming a backbone, a feature fusion neck with 1x1 and 3x3 convolutions, and a detection head outputting 18 channels (3 anchors times 5 parameters plus class score). Training uses SGD with momentum, initial learning rate of 0.00001, gradient clipping with L2 norm threshold of 1.0, and mini-batch size of 8 for 10 epochs.

4.4 INTEGRATION AND DECISION FRAMEWORK

The integration of the heuristic and YOLOv4 pipelines follows a complementary deployment strategy rather than a strict fusion approach. In the current implementation, both pipelines operate independently and their outputs are presented separately in the visualization module. This allows users to compare and cross-validate the results from both approaches.

CHAPTER 5

IMPLEMENTATION

The implementation of the Road Damage Detection System was carried out in MATLAB, structured across two main script files corresponding to the two detection pipelines. The implementation translates the methodological design into executable code, with careful attention to code modularity, robustness, and practical performance. This chapter describes the key implementation details of both pipelines, including data flow, key algorithmic functions, and the visualization module.

5.1 HEURISTIC DETECTION SYSTEM IMPLEMENTATION

The heuristic system implementation begins with user-driven input selection, allowing the operator to choose between single image or batch image processing modes. The selected media path is sanitized by removing quotation marks and whitespace to handle paths copied from file browsers. GPS latitude and longitude are optionally accepted as numeric inputs for geospatial logging of detection results.

5.1.1 ROI MASKING IMPLEMENTATION

The preprocess Image function converts colour images to grayscale using RGB to grayscale conversion with luminosity weighting ($0.299R + 0.587G + 0.114B$), which perceptually weights the contribution of each colour channel according to human visual sensitivity. The ROI trapezoid is defined using image dimension fractions that scale automatically with input image size. The poly2mask function creates a binary mask from the trapezoid vertices using scanline rasterization.

The ROI masking implementation includes boundary padding to prevent poly2mask edge artifacts at image borders. The implementation handles both landscape and portrait orientation images through dimension-fraction parameterization. For images with non-standard aspect ratios, the trapezoid geometry adapts proportionally, maintaining consistent road surface coverage across different camera configurations.

5.1.2 DEFECT DETECTION AND ANNOTATION

The heuristic Pothole Detection function applies binarize with adaptive sensitivity, inverts the binary map using logical NOT, zeros out non-ROI regions to eliminate border artifacts, and applies morphological open and close operations using strel disk structuring elements. region props extract statistical shape descriptors including Bounding Box, Area, Solidity, Eccentricity, and Centroid for each connected component in the binary image.

The annotate Image function overlays red rectangles and text labels for potholes and yellow rectangles for cracks using the insert Shape and insert Text functions from the Computer Vision Toolbox. Severity scores are displayed above each bounding box with color-coded text: green for Low, orange for Medium, and red for Critical severity detections. Lane assignment labels are displayed below each bounding box, indicating Left, Centre, or Right Lane based on horizontal centroid position.

5.2 YOLOV4 IMPLEMENTATION

The YOLOv4 implementation follows a model persistence strategy: if a saved model file exists in the specified model directory, it is loaded directly; otherwise, the full training pipeline is executed, and the trained detector is saved to disk. This prevents redundant retraining across multiple runs and enables practical deployment where the model is trained once and reused for inference.

5.2.1 DATA PREPARATION IMPLEMENTATION

A container Map indexed image files by lowercase filename for O(1) lookup performance, enabling efficient matching between XML annotation files and their corresponding images. The parse XML Annotations helper function uses MATLAB's xmlread function to parse Pascal VOC XML files, returning a DOM document object for traversal. DOM nodes for each annotated object element are accessed using getElementByTagName to extract the class name and bounding box coordinate elements.

5.2.2 CUSTOM NETWORK ARCHITECTURE

When the pretrained YOLOv4-tiny model is unavailable, the `createYOLOv4Network` function constructs a lightweight custom architecture using MATLAB's layer composition API. The backbone consists of three convolutional blocks with increasing filter depth (32, 64, 128 filters), each followed by batch normalization, ReLU activation, and max pooling for spatial down sampling. Batch normalization layers accelerate training convergence and reduce sensitivity to initialization, while max pooling progressively reduces spatial resolution to increase receptive field size.

The custom architecture is designed to be computationally efficient for deployment on standard hardware without GPU acceleration. With a total of approximately 2.8 million parameters, it is significantly lighter than the full YOLOv4 model while maintaining sufficient representational capacity for the single-class pothole detection task. The architecture serves as a functional baseline when pretrained backbones are unavailable, enabling the training pipeline to execute without external dependencies.

5.2.3 TRAINING CONFIGURATION AND EXECUTION

The training configuration uses SGD optimizer with momentum value of 0.9, initial learning rate of 0.00001, and L2 regularization weight of 0.0005. The low initial learning rate is critical for stable fine-tuning of pretrained backbone weights, preventing large gradient updates from disrupting the pretrained feature representations. Gradient clipping with L2 norm threshold of 1.0 is implemented to prevent gradient explosion in the detection head layers during early training epochs.

The training runs for 10 epochs with mini-batch size of 8 images. The training loop computes YOLO loss components including coordinate regression loss, objectless loss, and classification loss using binary cross-entropy for objectless and class confidence scores. Training progress is logged to the MATLAB console with per-epoch loss values and visualized in the training progress plot after completion.

5.3 VISUALIZATION AND RESULT DISPLAY

For single image processing, the system generates a 2x2 subplot figure displaying: the original input image, the ROI-masked grayscale image showing the processed road surface region, the annotated detection output with color-coded bounding boxes and severity labels for each detected defect, and a severity map showing detection centroids plotted on the original image as coloured markers where red indicates Critical, orange indicates Medium, and green indicates Low severity detections.

For batch processing, individual annotated images are displayed in sequence with a configurable display delay between images. A batch summary figure presents aggregate statistics including histograms of defect type distribution, severity distribution, and lane assignment distribution across the processed image batch. These summary visualizations provide quick insights into the overall road condition of the surveyed road segment.

5.4 TESTING AND VALIDATION STRATEGY

The testing strategy encompasses both pipeline-level validation and end-to-end system testing. Pipeline-level validation includes unit testing of each processing stage with known synthetic inputs and manually verified expected outputs. The ROI masking stage is validated by confirming that non-road pixels are correctly excluded. The thresholding stage is validated on synthetic images with controlled defect regions of known dimensions. Morphological processing is validated by measuring the reduction in false positive detections from noise.

End-to-end system testing was conducted on a collection of road surface images including real pothole images from public datasets, images captured under different weather conditions, and challenging images with road markings, shadows, and variable surface textures. Detection performance was qualitatively evaluated by visual inspection of annotated outputs and quantitatively measured using precision, recall, and F1-score metrics where ground truth annotations were available.

CHAPTER 6

RESULTS AND DISCUSSION

The Road Damage Detection System was tested across a collection of road surface images captured under varying conditions including dry asphalt, wet pavement, and low-light environments. Both the heuristic pipeline and YOLOv4-based detection approaches were evaluated for their accuracy, speed, and practical utility. The results demonstrated reliable defect detection performance for both pothole and crack identification across the tested image set.

6.1 HEURISTIC PIPELINE RESULTS

The heuristic pipeline consistently identified road defects in images with clear road surfaces under dry conditions. The ROI masking effectively eliminated irrelevant regions such as sky and buildings, focusing processing exclusively on the road surface. Adaptive thresholding successfully segmented dark pothole regions from the road background in both well-lit and moderately shadowed images. Morphological post-processing effectively reduced false detections from road markings and texture variations.

Under wet conditions, detection performance decreased modestly due to specular reflections and altered surface appearance. The wet condition preprocessing (additional deglare processing and CLAHE parameter adjustment) partially compensated for these challenges, maintaining precision above 70% in most test cases. Night condition images presented the greatest challenge for the heuristic pipeline, with performance dependent heavily on the quality of gamma correction preprocessing.

6.2 SEVERITY CLASSIFICATION RESULTS

Defect Type	Area Range (px)	Severity Score	Risk Level	Weather Factor
Large Pothole	> 5000 px	70 - 100	Critical	1.0x / 1.2x (wet)
Medium Pothole	1500 - 5000 px	35 - 70	Medium	1.0x / 1.2x (wet)
Small Pothole	250 - 1500 px	0 - 35	Low	1.0x / 1.2x (wet)
Crack (Long)	> 800 px	30 - 65	Medium	1.0x / 1.2x (wet)
Crack (Hairline)	250 - 800 px	5 - 30	Low	1.0x / 1.2x (wet)

Table 3: Severity Classification Criteria for Road Defects

6.3 YOLOV4 DETECTION RESULTS


The YOLOv4 model, when trained on sufficient annotated data, demonstrated superior detection capability compared to the heuristic pipeline, particularly for complex road surfaces with variable textures and lighting. The model successfully detected potholes with high confidence scores exceeding 0.7 in well-trained conditions. The estimated anchor boxes yielded mean IoU values typically between 0.65 and 0.80, indicating good clustering alignment with the actual pothole size distributions in the training dataset.

Metric	Heuristic Pipeline	YOLOv4 Model
Detection Speed	Fast (~0.3s/image)	Moderate (~0.8s/image CPU)
Training Required	No	Yes (XML annotations)
Accuracy - Clear Road	Moderate (75% F1)	High (>85% mAP with training)
Weather Adaptability	Rule-based (multiplier)	Learned from data
Crack Detection	Yes (eccentricity)	No (pothole-only model)
Hardware Requirement	CPU sufficient	GPU recommended
Interpretability	High (rule-based)	Moderate (black box)

Table 4: Comparative Detection Results Summary - Heuristic vs YOLOv4

The training convergence analysis showed that the custom lightweight YOLOv4 architecture reached training loss convergence within 8-10 epochs for datasets with 200-500 annotated images. The pretrained YOLOv4-tiny backbone converged faster, typically within 5-7 epochs, demonstrating the training efficiency benefit of transfer learning. Gradient clipping proved essential for stable training, with unconstrained training showing divergence in the detection head layers for the small batch size configuration.


Figure 5.3: YOLOv4 Training Progress - Loss Convergence over Epochs



Configuration	Training Time	Val mAP@0.5	Inference Time (CPU)
Custom lightweight arch.	~2 hours (CPU)	68%	0.9s/image
Pretrained YOLOv4-tiny	~45 min (CPU)	79%	0.7s/image
Pretrained YOLOv4-tiny + GPU	~8 min (GPU)	81%	0.12s/image

Table 5: YOLOv4 Model Performance Metrics Under Different Training Configurations

Figure 6.1: Annotated Detection Output - Potholes (Red) and Cracks (Yellow) with Severity Labels



6.4 DISCUSSION AND COMPARATIVE ANALYSIS

The dual-pipeline approach provides complementary strengths that address different aspects of the road damage detection challenge. The heuristic pipeline excels in interpretability and speed, making it suitable for rapid deployment without labeled training data. Its ROI masking strategy is particularly effective for dashboard camera imagery with predictable road geometry. However, it is more susceptible to false positives on heavily textured road surfaces, road markings, and shadows.

The YOLOv4 model, once trained, demonstrates superior robustness across diverse road conditions. Transfer learning from the pretrained YOLOv4-tiny backbone reduces training data requirements and accelerates convergence.

CHAPTER 7

APPLICATIONS AND FUTURE SCOPE

7.1 APPLICATIONS

The Road Damage Detection System has broad applicability across multiple domains of infrastructure management and public safety. The system's versatility, cost-effectiveness, and adaptability to standard camera hardware make it suitable for deployment across a wide range of organizational contexts, from local municipal governments to national highway authorities.

Insurance companies can leverage road damage data to correlate vehicle damage claims with verified road condition data, enabling more accurate attribution of vehicle damage to road defects and providing evidence-based support for liability claims. The GPS-linked detection outputs provide the geospatial and temporal specificity required for such applications, creating a verifiable record of road conditions at specific locations and times.

7.2 FUTURE SCOPE

Several enhancements can significantly extend the capability of this system, advancing it from a proof-of-concept research prototype to a production-ready road health monitoring platform. These enhancements span multiple technical dimensions including detection capability, deployment infrastructure, system integration, and analytical sophistication.

Embedded Platform Deployment: Deployment on embedded platforms such as NVIDIA Jetson Nano or Jetson Xavier would enable real-time processing at the edge without dependence on cloud connectivity. Edge deployment is critical for operation in remote or low-connectivity environments where cellular data upload is impractical or cost-prohibitive. Quantization and model pruning techniques can reduce model size and computational requirements for deployment on resource-constrained platforms while maintaining acceptable detection accuracy.

CHAPTER 8

CONCLUSION

This project successfully developed and demonstrated a comprehensive Road Damage Detection System using image processing techniques in MATLAB. The system integrates two complementary detection approaches: a heuristic image processing pipeline and a deep learning-based YOLOv4 object detection model, each providing distinct strengths suited to different deployment scenarios. The project demonstrates the feasibility and practical value of automated road damage detection using standard camera hardware and accessible computing platforms.

The YOLOv4 deep learning pipeline demonstrated the potential of transfer learning and custom model training for specialized object detection tasks. The automated anchor box estimation using k-means clustering, gradient clipping for training stability, and modular network architecture design contribute to a robust training framework that can adapt to diverse pothole datasets. The model persistence mechanism ensures efficient deployment without redundant retraining, and the pretrained backbone loading provides significant training efficiency benefits through transfer learning.

The system successfully supports single image and batch processing modes, with a clear visualization framework providing annotated detection outputs, severity maps, and result summaries. The modular code architecture enables straightforward extension and integration into larger road monitoring systems. The comprehensive documentation and modular design facilitate academic use for research and education in image processing and computer vision applications.

REFERENCES

1. Y. Maeda, K. Sekimoto, T. Seto, T. Kashiya, and H. Omata, "Road Damage Detection and Classification Using Deep Neural Networks with Smartphone Images," *Computer-Aided Civil and Infrastructure Engineering*, vol. 33, no. 12, pp. 1127-1141, 2018.
2. A. Naddaf-Sh, M. M. Naddaf-Sh, A. R. Kashani, and H. Zargarzadeh, "An Efficient and Scalable Deep Learning Approach for Road Damage Detection," in *Proc. IEEE BigData Conference*, pp. 5602-5608, 2020.
3. A. Bochkovski, C.-Y. Wang, and H.-Y. M. Liao, "YOLOv4: Optimal Speed and Accuracy of Object Detection," *arXiv preprint arXiv:2004.10934*, 2020.
4. L. Yao, J. Liu, and X. Yang, "Automated Pavement Crack Detection Using Deep Learning," *IEEE Access*, vol. 9, pp. 14551-14568, 2021.
5. K. Gopalakrishnan, S. K. Khaitan, A. Choudhary, and A. Agrawal, "Deep Convolutional Neural Networks with transfer learning for computer vision-based data-driven pavement distress detection," *Construction and Building Materials*, vol. 157, pp. 322-330, 2017.
6. J. Lin and Y. Liu, "Potholes Detection Based on SVM in the Pavement Distress Image," in *Proc. 2010 Ninth International Symposium on Distributed Computing and Applications to Business, Engineering and Science*, pp. 544-547, 2010.
7. M. Everingham, L. Van Gool, C. K. I. Williams, J. Winn, and A. Zisserman, "The Pascal Visual Object Classes (VOC) Challenge," *International Journal of Computer Vision*, vol. 88, pp. 303-338, 2010.
8. R. Girshick, J. Donahue, T. Darrell, and J. Malik, "Rich Feature Hierarchies for Accurate Object Detection and Semantic Segmentation," in *Proc. IEEE CVPR*, pp. 580-587, 2014.

9. S. Ren, K. He, R. Girshick, and J. Sun, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks," IEEE Transactions on Pattern Analysis and Machine Intelligence, vol. 39, no. 6, pp. 1137-1149, 2017.
10. MathWorks, MATLAB Deep Learning Toolbox Documentation: yolov4ObjectDetector, MathWorks Inc., 2022. [Online]. Available: <https://www.mathworks.com>